Multi-Agent Reinforcement Learning for Cooperative Navigation

Diana Huang [1]   Shalini Keshavamurthy [1]   Nitin Viswanathan [1]

Abstract

Multi-agent scenarios frequently come up in the real world but traditional reinforcement learning algorithms do not perform well on them due to constantly changing environments from the perspective of any one agent. We replicate multi-agent deep deterministic policy gradients (MAD-DPG), an algorithm tailored to multi-agent scenarios, and evaluate its performance in a 3-agent cooperative navigation OpenAI environment. Additionally, we perform many experiments to explore how changing hyperparameters and action exploration strategies affects performance. Finally, we enhance MADDPG to address its main weakness, which is that it does not scale well to larger numbers of agents.

Our results confirm that MADDPG performs significantly better than a single-agent policy gradient approach. Additionally, we show that our strategy to allow MADDPG to better scale to larger numbers of agents does not cause it to perform significantly worse, indicating that it could be a better strategy than the vanilla algorithm.

1. Introduction

There are many applications where multiple agents need to learn how to act together, such as multiplayer games (Peng et al., 2017), multi-robot control (Matignon et al., 2012a), communication scenarios (Foerster et al., 2016), and even self-play (Sukhbaatar et al., 2017).

Traditional reinforcement learning approaches focus on training a single agent and as a result they do not work well on multi-agent problems because the environment for any single agent changes over time as other agents change their policies, leading to instability during training and high variance in results (Matignon et al., 2012b). Additionally,

---
*Equal contribution [1]Stanford University, Palo Alto, USA. Correspondence to: Diana Huang <hxydiana@stanford.edu>, Shalini Keshavamurthy <skmurthy@stanford.edu>, Nitin Viswanathan <nitin.viswanathan@gmail.com>.

experience replay as used with Q-learning cannot be directly applied in multi-agent scenarios, significantly hampering stable learning and data efficiency.

Research into multi-agent reinforcement learning has attempted to develop new approaches specifically for multi-agent scenarios. One such recently developed algorithm is multi-agent deep deterministic policy gradients algorithm (MADDPG), which obtains significantly improved results over other approaches by modifying DDPG to train agents while incorporating information from other agents (Lowe et al., 2017).

Our paper has 3 key contributions:

- We replicate the MADDPG algorithm and apply it to the OpenAI cooperative navigation environment
- We experiment with various hyperparameters and action exploration strategies to show how MADDPG's performance is affected by them
- We extend MADDPG to handle larger numbers of agents efficiently and explore how performance changes with the modified algorithm

2. Related Work

2.1. Single-agent methods

Prior work on multi-agent scenarios has attempted to use single-agent methods that train agents independently with limited success. Prior work has applied Q-learning to cooperative multi-agent systems but found it to be less robust than in single-agent settings (Sandholm & Crites, 1996), (Claus & Boutilier, 1998). Other work has attempted to modify Q-learning to better fit multi-agent scenarios, but focused primarily on cooperative environments instead of also considering competitive ones (Lauer & Riedmiller, 2000).

Other work has tried to address the problem by using incremental learning and policy gradient approaches (Buffet et al., 2007). These approaches struggle because from the point of view of any one agent, the environment is non-stationary as other agents change their policies (Busoniu et al., 2008), (Buşoniu et al., 2010). We expand on these issues in the Approach section.

2.2. Multi-agent methods

Less prior work exists specifically attempting to tailor approaches for multi-agent scenarios, and even less with experimental results. (Boutilier, 1996) discusses coordination and planning between multiple agents using multi-agent markov decision processes where all the agents know information about other agents. (Chalkiadakis & Boutilier, 2003) proposes the use of bayesian models for optimal exploration in multi-agent scenarios.

Recent work has proposed multi-agent deep deterministic policy gradients (MADDPG), an approach based on DDPG that is tailored to multi-agent scenarios (Lowe et al., 2017). MADDPG combines concepts from (Sutton & Barto, 1998) and (Lillicrap et al., 2015) to train agents using information from other agents via a centralized critic function. MADDPG significantly outperforms single-agent approaches and is a generalizable method that can work in a variety of cooperative and competitivee environments, whereas most prior work focuses only on cooperative environments. We choose to based our work on MADDPG.

Other work has explored using a centralized critic function like MADDPG does during training in a Starcraft micromanagement environment (Foerster et al., 2017).

3. Approach

3.1. Environment

We use the cooperative navigation OpenAI Gym environment for our experiments (OpenAI, 2017). This is a 2-D world with N moving agents and N fixed landmarks ($N = 3$ by default) where the agents must learn to each move to a different landmark while avoiding collisions with each other.

Every agent has its own state that is an 18-dimensional vector containing its own position and velocity, relative distance to each of the landmarks, and relative distance to the other agents. The action that every agent can take is represented as a 5-dimensional vector, with 4 dimensions representing movement up/down/left/right and the last dimension representing a no-move action. This is a continuous action space with the values for each action dimension representing acceleration in a given direction.

At each timestep, all agents are rewarded based on the distance between landmarks and agents; the closer the agents are to separate landmarks, the higher the reward. If every agent was exactly on top of a different landmark, the reward at that timestep would be 0. Agents have a non-zero size so it is possible for them to collide in this environment. Agents receive an additional penalty of -1 if there are collisions at any timestep. Therefore, in order to maximize reward the agents need to each learn to move to a different landmark


Figure 1. Cooperative navigation environment. Agents must learn to each navigate to a different landmark without colliding.

while also avoiding collisions with each other.

We define an episode to be 25 timesteps in length as there is no defined terminal state in this environment.

3.2. Single-agent Policy Gradient

We implement a vanilla single-agent policy gradient algorithm as a baseline. Policy gradient algorithms seek to determine the optimal policy by directly adjusting the parameters θ of a policy π_θ in order to maximize the objective function $J(\theta) = \mathbb{E}_{\pi_\theta}[R]$. We also implemented a baseline and advantage for the policy gradients to help reduce variance and increase stability.

Letting $\hat{A} = R_t - b(s_t)$ be the advantage estimate at time t, we can write the gradient of the policy as:

$$\nabla_\theta J(\theta) = \mathbb{E}_{s,\pi_\theta}[\nabla_\theta log \pi_\theta(a_t|s_t)\hat{A}_t]$$

In the above equation, the expectation is taken over many sample trajectories through the environment using π_θ. The main downside of single-agent policy gradient algorithms is that they are already susceptible to high variance and executing them in a multi-agent environments only amplifies the problem. From a single agent's perspective not only does the world change as other agents move, but the policies other agents follow change as well. Additionally, an agent might not get the correct gradient signal if for example it takes a good action but other agents take bad actions, leading to overall decrease in reward. In this case, the policy gradient could cause the agent to move away from taking the good action.

3.3. MADDPG

We recreate MADDPG as our main algorithm for the cooperative navigation task (Lowe et al., 2017). MADDPG takes DDPG and tailors it to multi-agent environmnnts, resulting in significantly better performance than single-agent methods. MADDPG trains separate policy networks for every agent, but uses a centralized action-value function that
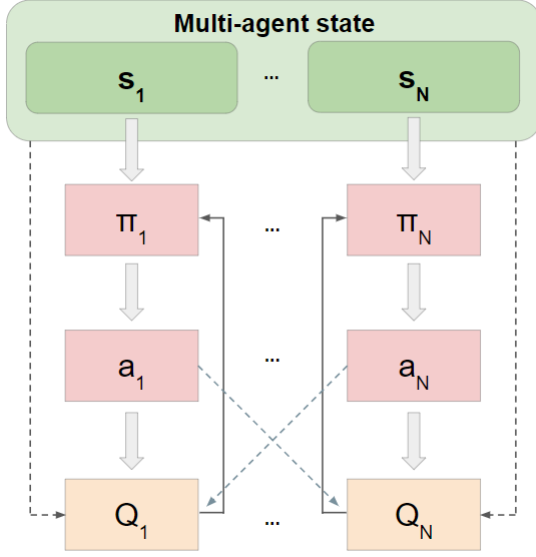
Figure 2. MADDPG algorithm with centralized Q functions. As indicating by the dotted lines, note that the Q functions take the entire multi-agent state as an input in addition to the actions of all agents.

incorporates information across all agents to ensure a more stable learning process.

Consider a game with N agents, and let $\pi = \{\pi_1, ..., \pi_N\}$ be the associated set of agent policies that are parametrized by $\theta_1, ..., \theta_N$. Additionally let $s = \{s_1, ..., s_N\}$ represent the set of individual agent observations. When using MADDPG to train agents, every agent additionally maintains a critic function that outputs a Q-value given the states and actions of all agents (see Figure 2):

$$Q_i^\pi(s, a_1, ..., a_N)$$

Standard Q-learning techniques are used to update the critic function. Note that for the cooperative navigation environment that we use for our experiments, the Q-values will be the same for all agents as the reward for all agents is the same. However, this Q-function formulation where every agent learns its own Q_i^π allows MADDPG to generalize to environments where agents have different reward functions (e.g. adversarial games).

During training, agents now have information not just about the locations of other agents, but about the actions they will take as well. Therefore, the environment is stationary for any agent as it does not depend on other agent policies:

$$P(s_i'|s_i, a_1, ..., a_N, \pi_1, ..., \pi_N) = P(s_i'|s_i, a_1, ..., a_N)$$

The policy network updates for an agent are determined by maximizing the agents centralized Q-values. Therefore, gradient update is similar to the update used in the vanilla

policy gradient except that it uses the centralized Q-value instead of only looking at the returns across trajectories:

$$\nabla_{\theta_i} J(\theta_i) = \mathbb{E}_{s,\pi_i}[\nabla_{\theta_i} log\pi_i(a_i|s_i)Q_i^\pi(s, a_1, ..., a_N)]$$

By incorporating the actions of all agents into the action-value function during training instead of only the action of an individual agent, MADDPG produces better and more stable results than traditional policy gradient or Q-learning approaches. An additional advantage of MADDPG is that it supports experience replay. When training, we apply actions from agent policy networks and store the result state transition in an experience replay buffer. Periodically, we update network parameters by sampling training data from the replay buffer and applying MADDPG to the sampled transitions. We also make use of target policy networks and target Q-networks to further stabilize training.

During policy execution, agents only use their learned policy networks and therefore do not use any information about other agents; MADDPG is an enhancement to the training process while leaving the execution phase unchanged.

3.4. Scaling MADDPG by only use neighboring agents

In addition to implementing MADDPG as described in the previous section, we modify it to scale better for supporting large numbers of agents. The Q function described above in MADDPG algorithm increases linearly in input size as more agents are added to the scenario, as it has to take every agent's observation and action as an input. In order to reduce the additional computation time needed as N grows, we modify MADDPG to only look at the K nearest neighbors during training, where $K < N - 1$. The mathematical formulation of the algorithm remains the same, except that the Q_i^π only takes in the states and actions of agent i and its nearest K neighbors.

3.5. Ornstein-Uhlenbeck action noise

One of the exploration methods we have tried in our experiments is the Ornstein-Uhlenbeck (OU) noise process (Bibbona et al., 2008). OU noise is a stochastic process that roughly models Brownian motion and is also known as the correlated additive Gaussian action space noise. The idea behind using this process is that if the noise generated at a given time step is correlated to previous noise, it will tend to stay in the same direction for longer durations instead of immediately canceling itself out, which will consequently allow the agent to go towards the goal faster. So in the case of continuous actions it is better to have temporally correlated exploration that suit the dynamics to get smoother transitions. This noise process is popularly used with DDPG, so we implement it as well to see if MADDPG performs well with it.

4. Experiments and Results

4.1. Evaluation Metrics

In order to evaluate the performance of our algorithms, every 250 training steps we do a evaluation run of 40 episodes using the latest learned policies and calculate various metrics over them. Data from evaluation runs do not go into training set. Our primary evaluation metric for all experiments is average reward, which corresponds to the distance between agents and the target landmarks while also accounting for collisions that occur along the way. In addition to average reward, we also look at average distance between landmarks and agents, for more intuitive interpretation of the performance. Here we provide a detailed definition of average reward and average distance used in this paper, unless specified otherwise.

Average reward. Average reward is calculated by summing up rewards received from all agents, all timesteps within an episode, and then averaging across multiple episodes. The range of average reward is zero to negative infinity.

Average distance. Average distance is the sum of distances between each agent and its closest landmark, across all agents and all timesteps within an episode, and then averaged across episodes.

Besides the metrics defined above, each experiment may additional context-specific evaluation metric, which will be defined for that experiment.

4.2. Implementation details

Table 1 lists the default hyperparameters used in our experiments. We use fully connected neural networks with a softmax activation at the final output layer to represent the policy networks and use fully connected network without output layer activation for the Q networks.

Episode Length. The number of of actions every agent can take from the start state until the episode ends.

Batch size. The number of sequences of agent observations, actions, rewards and next observations used to train agents. In the case of naive PG, the sequences are in the order they happen, whereas with experience replay buffer, the sequences are randomly drawn and independent of each other.

Training frequency. The frequency of how often we sample a training batch from the experience replay buffer (which stores 10^6 timesteps) and update network parameters with it, in timesteps. At timesteps where we are not training we use the current agent policies to run episodes and obtain new training data to store in the replay buffer.

Gamma. The discount factor which models the future rewards as a fraction of the immediate rewards. It quantifies

how much importance we give for future rewards.

Learning rate. Specifies the magnitude of step that is taken towards the solution.

Tau. Affects how fast the target policy and Q networks update.

Hyperparameter	Value
# of hidden layers	2
# of units per hidden layer	128
Episode length (timesteps)	25
Training batch size	1250
# of training batches	60,000
Training frequency (timesteps)	100
Gamma (γ)	0.95
Learning Rate (α)	0.01
Tau (τ)	0.01

Table 1. Default hyperparameters used for training MADDPG.

4.3. Experiments

Baseline comparison. As a baseline, we compare the performance of our single-agent policy gradient baseline and MADDPG with the baseline hyperparameters as given in Table 1. As shown in figures 3 and 4, the policy gradient approach has difficulties learning a good policy. It fails to both increase reward and decrease distance to landmarks while training, and produces extremely noisy results. In contrast, MADDPG's performance across both metrics steadily improves as training progresses and it exhibits much less variation in results, indicating that training is more stable. We believe MADDPG's good performance comes from the fact that it uses a centralized critic function and that we can use experience replay with it.

Batch size comparison. Next, we examine the effect of training batch size on MADDPG performance. In addition to the default of 1,250 timesteps, we also explored using 6,250 and 25,600 timesteps per training batch. Our hypothesis was that increasing training batch size (but still performing the same number of training batches) would improve performance, as incorporating more training timesteps into each update makes it more likely that the resulting gradient is in the right direction. Figure 5 shows our results and it does not appear that increasing batch size results in better performance in our runs. As a result of this experiment, we use a batch size of 1,250 timesteps for our other experiments as it takes the least amount of time to run.

Action exploration strategies. Action exploration and balancing it with exploitation is a critical part of determining optimal parameters. (Lowe et al., 2017) performs action exploration by adding noise to the unnormalized policy network output when sampling actions and then applying a

Figure 3. Average episode reward achieved by our naive PG and MADDPG implementations during training.


*Figure 4.* Average distance from landmarks with Naive PG algorithm and MADDPG algorithm.

Figure 5. Average episode reward during training with different training batch sizes.
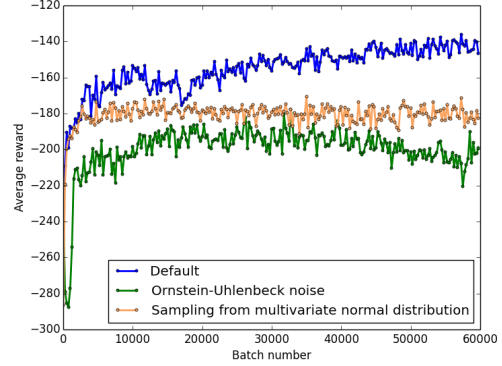

Figure 6. Average reward gained by using different exploration strategies.

softmax function to obtain the final action. This approach seems unintuitive to us and the rational for it is not mentioned in the paper, so we experiment with this and other approaches to better understand performance.

In addition to implementing the MADDPG action exploration methodology, we also try applying an Ornstein-Uhlenbeck process with $\sigma = 0.3, \theta = 0.15$ and sampling from multivariate normal distribution with a trainable standard deviation. Figure 6 shows the average rewards with different action exploration strategies. We see that the original MADDPG action exploration methodology does best, supporting the findings in (Plappert et al., 2017) that adding noise directly to the network improves performance. One possible reason for is this is that because the reward function is well-shaped and reward is given at every timestep, not as much exploration of the action space is needed.

We additionally look at the standard deviation of the deviation from the network output caused by the various exploration techniques in figure 7. We see that the default approach from (Lowe et al., 2017) has the lowest deviation from the network outputs, supporting the theory that not as much exploration is needed. The other two methods we try both perform more exploration and have higher standard deviations, likely leading to lower resulting reward because limited episode length does not allow much room for mistakes and corrections.

Parameter space noise for exploration. Previously, we considered exploration by adding noise to the output action. Then we experimented with adding noise to the parameters of the actor network. (Plappert et al., 2017) presents the idea of adding noise directly to the agent's parameters. The reasoning behind this is as follows. In the continuous action space, actions are sampled according to a stochastic policy. For a fixed state s, we will most certainly obtain a different action when that state is sampled again since action space

Figure 7. Standard deviation of effective euclidean distance for different exploration strategies.


Figure 8. Average reward with and without network parameter noise


Figure 9. Average reward gained by each agent in environment with different number of agents and landmarks

noise is independent of the current state s. On the other hand, if the network parameters are perturbed at the beginning of each episode, the same action will be taken for the same state s. This ensures dependency of actions and state, eventually leading to consistency of outputs. We have used layer normalization between hidden layers, so that the same perturbation scale can be used across all layers. We have used $\sigma = 0.2$ as the scale for parameter space noise. Figure 8 shows the default case with default noise and the default case with network parameter noise. As we can see in the plot, the average rewards are much higher with parameter noise compared to the default results.

Changing network size. We experiment with reducing the size of neural network layers for the policy and Q networks to see the impact on performance. We observe that the best achievable reward in a training run with 2-layer 64-unit networks (as opposed to the default of 2 layers with 128-unit networks) is around -175, which is noticeably lower than the -150 achieved with the default network size. This shows that increased complexity of neural networks does help to define the right actions in the cooperative navigation task.

Increased number of agents and landmarks. Figure 9 show the results from training different number of agents using MADDPG. For each result, the number of agents is equal to the number of landmarks so every agent is still trying to navigate to a different landmark. Here we use average reward per agent to measure performance, which is average reward defined earlier divided by number of agents. As the number of agents increases, we can see that the average episode reward per agent decreases. This shows that training becomes more difficult as the environment space becomes more complex.

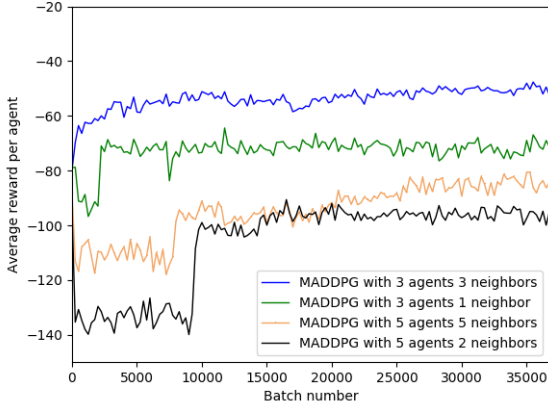**Scaling MADDPG by only using neighboring agents.**

Figure 10. Average reward gained by each agent with different number of neighboring agents information

Figure 10 shows the performance of our modified MAD-DPG implementation that only uses information from the K nearest neighbors during training, where $K < N - 1$. We evaluate our implementation in two different scenarios, one with 3 agents but only using information from 1 neighbor instead of both, and one with 5 agents and using information from 2 neighbors instead of all 4. We compare the performance in these to scenarios with 3 and 5 agents using information from all other agents. The modified MADDPG algorithm performs slightly worse at both 3 and 5 agents, which is expected as less information is used during training. However, we can see that the gap in performance at $N = 5$ is smaller than at $N = 3$, suggesting the possibility that as N increases further, the gap might continue to remain small.

This hypothesis requires more runs to validate (e.g. running a scenario with 7 agents with and without limited information on neighbors), but suggests that having information on nearest neighbors is the most relevant information to avoiding collisions. As long as an agent moves towards a landmark while also avoiding collisions with its nearest neighbors, it is likely going to increase overall episode reward. Another key advantage of our modified MADDPG implementation is that the Q network has fewer parameters, speeding up training especially as N continues to increase relative to K.

5. Conclusion

We introduce the problem of multi-agent reinforcement learning and show why traditional single-agent approaches are ineffective on these problems. We implement a single-agent policy gradient method as a baseline and MADDPG, an algorithm designed for multi-agent scenarios, and apply them to the OpenAI cooperative navigation environment.

We perform a variety of experiments involving changing hyperparameters, environment settings, and more and analyze the resulting performance. Finally, we modify MADDPG to look at a fixed number of nearest neighbors during training to enable it to scale better to larger numbers of agents.

We show that MADDPG clearly outperforms policy gradient approaches as it is able to steadily improvement performance over training batches while policy gradient fails to improve. Our experiments also show that not much noise on policy network outputs is needed for action exploration in the cooperative navigation scenario. And using parameter space noise exploration can be a better choice for training the agent policy. Increasing batch size beyond 1,250 timesteps does not appear to noticeably impact training performance, while decreasing the policy network and Q-network from 2 layers and 128 units to 2 layers and 64 units does negatively impact training performance.

As MADDPG is run with more agents, we see that the average reward per agent decreases as a result of the environment becoming more complicated. Running vanilla MADDPG on larger numbers of agents significantly increases computation time due to the increased size of the Q-function's input space, but our modification of MADDPG that only looks at a set number of nearby neighbors does not have this same problem and also performs close to the default runs, indicating that generalizing MADDPG could be a viable way to scale it to support scenarios with more agents.

6. Future Work

We would like to try our approach on other multi-agent environments, in particular adversarial ones, to compare how MADDPG performs vs. policy gradients. It would also be interesting to explore how performance changes if agents are trained using different policies - e.g. in an adversarial game, what would happen if one set of agents were trained with MADDPG and a competing set were trained with policy gradients.

Our results from generalizing MADDPG to only look at N nearest neighbors show that an agent might not even need to know about all other agents during training to achieve good results. We would like to further experiment with generalization, in particular by running against even larger numbers of agents.

Additionally, we would like to explore alternative forms of generalization. Instead of Q_i^π looking at the K nearest neighbors to agent i for some fixed K, the Q-function could instead look at all neighbors within a particular distance to agent i. This way, the agents would have more information available to them when it is particularly important (the closer agents are to each other, the more likely it is to have a collision) while minimizing input space size when

information about additional agents is less necessary.

Contributions

Diana wrote the initial implementation of MADDPG, Shalini wrote our logging/plotting code, and Nitin wrote the initial implementation of the policy gradient baseline. All of us iterated on the code with bugfixes and enhancements and ran experiments using our implementations.

References

Bibbona, Enrico, Panfilo, Gianna, and Tavella, Patrizia. The ornsteinuhlenbeck process as a model of a low pass filtered white noise. *Metrologia*, 45 (6):S117, 2008. URL http://stacks.iop.org/0026-1394/45/i=6/a=S17.

Boutilier, Craig. Planning, learning and coordination in multiagent decision processes. In *Proceedings of the 6th conference on Theoretical aspects of rationality and knowledge*, pp. 195–210. Morgan Kaufmann Publishers Inc., 1996.

Buffet, Olivier, Dutech, Alain, and Charpillet, François. Shaping multi-agent systems with gradient reinforcement learning. *Autonomous Agents and Multi-Agent Systems*, 15(2):197–220, Oct 2007. ISSN 1573-7454. doi: 10.1007/s10458-006-9010-5. URL https://doi.org/10.1007/s10458-006-9010-5.

Busoniu, Lucian, Babuska, Robert, and De Schutter, Bart. A comprehensive survey of multiagent reinforcement learning. *IEEE Trans. Systems, Man, and Cybernetics, Part C*, 38(2):156–172, 2008.

Buşoniu, Lucian, Babuška, Robert, and De Schutter, Bart. Multi-agent reinforcement learning: An overview. In *Innovations in multi-agent systems and applications-1*, pp. 183–221. Springer, 2010.

Chalkiadakis, Georgios and Boutilier, Craig. Coordination in multiagent reinforcement learning: A bayesian approach. In *Proceedings of the second international joint conference on Autonomous agents and multiagent systems*, pp. 709–716. ACM, 2003.

Claus, Caroline and Boutilier, Craig. The dynamics of reinforcement learning in cooperative multiagent systems, 1998.

Foerster, J. N., Assael, Y. M., de Freitas, N., and Whiteson, S. Learning to communicate with deep multi-agent reinforcement learning, 2016.

Foerster, Jakob, Farquhar, Gregory, Afouras, Triantafyllos, Nardelli, Nantas, and Whiteson, Shimon. Counterfactual multi-agent policy gradients. *arXiv preprint arXiv:1705.08926*, 2017.

Lauer, M. and Riedmiller, M. An algorithm for distributed reinforcement learning in cooperative multi-agent systems, 2000.

Lillicrap, Timothy P, Hunt, Jonathan J, Pritzel, Alexander, Heess, Nicolas, Erez, Tom, Tassa, Yuval, Silver, David, and Wierstra, Daan. Continuous control with deep reinforcement learning. *arXiv preprint arXiv:1509.02971*, 2015.

Lowe, Ryan, Wu, Yi, Tamar, Aviv, Harb, Jean, Abbeel, Pieter, and Mordatch, Igor. Multi-agent actor-critic for mixed cooperative-competitive environments. *CoRR*, abs/1706.02275, 2017. URL http://arxiv.org/abs/1706.02275.

Matignon, L., Jeanpierre, L., and et al., A.-I. Mouaddib. Coordinated multi-robot exploration under communication constraints using decentralized markov decision processes, 2012a.

Matignon, L., Laurent, G. J., and Fort-Piat, N. Le. Independent reinforcement learners in cooperative markov games: a survey regarding coordination problems, 2012b.

OpenAI. Multi-agent particle environments. https://github.com/openai/multiagent-particle-envs, 2017.

Peng, P., Yuan, Q., Wen, Y., Yang, Y., Tang, Z., Long, H., and Wang, J. Multiagent bidirectionallycoordinated nets for learning to play starcraft combat games, 2017.

Plappert, Matthias, Houthooft, Rein, Dhariwal, Prafulla, Sidor, Szymon, Chen, Richard Y, Chen, Xi, Asfour, Tamim, Abbeel, Pieter, and Andrychowicz, Marcin. Parameter space noise for exploration. *arXiv preprint arXiv:1706.01905*, 2017.

Sandholm, Tuomas W. and Crites, Robert H. Multiagent reinforcement learning in the iterated prisoner's dilemma. *Biosystems*, 37(1):147 – 166, 1996. ISSN 0303-2647. doi: https://doi.org/10.1016/0303-2647(95)01551-5. URL http://www.sciencedirect.com/science/article/pii/0303264795015515.

Sukhbaatar, S., Kostrikov, I., Szlam, A., , and Fergus, R. Intrinsic motivation and automatic curricula via asymmetric self-play, 2017.

Sutton, R. S. and Barto, A. G. *Reinforcement Learning: An Introduction*. MIT press Cambridge, volume 1 edition, 1998.