
Generalized Multi-Agent Reinforcement Learning in Cooperative and Competitive Environments

Diana Huang^{*1} Shalini Keshavamurthy² Nitin Viswanathan^{*12}

1. Introduction

There are many applications where multiple agents need to learn how to act together, such as multiplayer games (cite), multi-robot control (cite), communication scenarios (cite), and even self-play (cite). Existing efforts in Multi-Agent Reinforcement Learning (MARL) have looked at both existing reinforcement learning methodologies as well as modifying them to better support multiple agents.

Traditional reinforcement learning approaches, focused on training a single agent, do not work well on these problems because the environment for any single agent changes over time as other agents change their policies, leading to instability during training and high variance in results (cite). Furthermore, experience replay as designed for traditional Q-learning cannot be applied in multi-agent scenarios, significantly hampering stable learning.

Recent work has achieved better results in multi-agent scenarios by modifying existing RL algorithms to specifically operate in this space, creating Multi-agent Deep Deterministic Policy Gradients (MADDPG) (Lowe et al., 2017). However, one key limitation of these approaches is that they do not generalize well to larger numbers of agents due to the growing computational complexity of incorporating information from all agents in training. Given the extensive training time necessary for even smaller scenarios, it is not feasible to retrain new networks for different possible numbers of agents.

We develop a generalized training approach for multi-agent scenarios to find policies that scale well when applied in environments with varying numbers of agents. We do so by modifying the training process of MADDPG to only train agents using information from nearby agents instead of information from every agent, allowing agents to train only based off of the location of a fixed number of nearby agents instead of every agent in the environment.

2. Approach

2.1. Environment

We use an OpenAI Gym environment with a variety of cooperative and competitive scenarios, as well as environments




Figure 1. Example environment initialization. Blue circles are agents and black dots are target landmarks.

that allow communication between agents (OpenAI, 2017). For our initial explorations, we focused on the cooperative navigation environment where there are N agents that can move in a 2-D world and N fixed landmarks. All agents are rewarded based on the collective distance between agents and landmarks, so in order to maximize reward the agents need to learn to each move on top of a different landmark. However, agents have a non-zero size and are penalized for colliding with each other. The agents need to learn to move to different landmarks without running into each other.

The observation space for the cooperative navigation environments consists of information about every agent’s position and velocity, relative distance to all of the landmarks, and relative distance to every other agent. An agent acts by choosing to move either up, down, left, right, or taking a no-move action.

As part of our post-milestone investigations, we plan to apply our methods to other multi-agent environments as well as they should generalize.

2.2. Methods

As an initial baseline, we have implemented the REINFORCE algorithm for our multi-agent scenario. With this methodology, every agent trains its own independent policy network based on the observed environment and possible actions it can take. We did not expect REINFORCE to perform well on the environment due to how much it changes across time steps - not only does the environment change

as agents move, but the policies they follow change as well. Additionally, individual agents might not always get the correct gradient signal as an agent could take a good action but if other agents took bad actions, leading to overall decrease in reward, the gradient for the agent taking the correct action could still be negative. Nevertheless, REINFORCE serves as a baseline for comparing against more sophisticated approaches.

We are currently in the process of implementing MADDPG from (Lowe et al., 2017) as a second baseline. MADDPG takes DDPG and tailors it to multi-agent scenarios, resulting in significantly better performance than single-agent methods such as REINFORCE and DDPG. MADDPG trains separate policy networks for every agent, but uses a centralized critic function that incorporates information across all agents to ensure a more stable learning process:

$$Q_i = 0$$

Our third algorithm, will be a modified version of MADDPG that trains agents in a generalized way. More specifically, we will still maintain a centralized critic function for

2.3. Experiments

For the milestone, we have been focusing on the cooperative navigation environment as mentioned previously. Our primary evaluation metric so far has been average reward, which in this environment corresponds to the distance between agents and the target landmarks while also accounting for collisions that occur along the way. In addition to average reward, in the future we will also look at success rate, where we define success as agents successfully reaching the target landmarks with 0 collisions, no matter how long it took for this to occur. This way, we are evaluating our methodology both based on how often it converges quickly and how often it converges without errors.

We will evaluate our model performance both during train and test time. During train time, we will see how quickly we are able to reach convergence and how stable the training process is across the different implemented methods. During test time, we will compare our different methods with the average reward over many trials as well as the variance of the average, and also see how performance changes as we change the number of agents present to test generalization. The original authors of MADDPG only tested scenarios with a maximum of 3 agents, so we will try increasing this to 4-6 agents and seeing how performance changes.

We expect to be able to match the results from MADDPG when evaluating on up to 3 agents. However, we aim to outperform MADDPG when applying the policies generated from training on 3 agents to scenarios with more than 3




Figure 2. Average reward over time when training PG with 1 layer and 16 hidden units




Figure 3. Average reward over time when training PG with 2 layers and 128 hidden units

agents, as we will be specifically tailoring our algorithm for generalization to larger numbers of agents.

TODO - insert table showing the parameters we used for the different policies.

2.4. Results

Figures 2.4 and 2.4 show our initial results from training our network. We can see that

TODO - write about the results. That we didn't get much improvement even with more network depth. 2 layers took a bit longer to train / reach the maximum, but

References

Lowe, Ryan, Wu, Yi, Tamar, Aviv, Harb, Jean, Abbeel, Pieter, and Mordatch, Igor. Multi-agent actor-critic for mixed cooperative-competitive environments. *CoRR*, abs/1706.02275, 2017. URL <http://arxiv.org/abs/1706.02275>.

OpenAI. Multi-agent particle environments. <https://github.com/openai/multiagent-particle-envs>, 2017.